# gem5's Standard Library

# Why a Standard Library?

When done without the library you must define *every part* of your simulation: Every Simobject, connected correctly to every port, for every part, no matter now small.
This can result in scripts of hundreds of lines of code even for the most basic of simulations.

This resulted in:

- A lot of duplicated code.
- Error-prone configurations.
- A lock of portability between different simulations setups.

In addition, while there is no "one size fits all" for gem5 users, most users have similar needs and requirements for their simulations; requiring only a few modifications off some commonly used configuration systems.
Prior to the creation of the standard library users would regularly circulte long complex scripts and hack at them endlessly.
Such practices inspired the creation of the gem5 Standard Library.

# What is the Standard Library?

The purpose of the gem5 Standard library is to provide a set of pre-defined components that can be used to build a simulation that does the majority of the work for you.

The the remainder not supported by the standard library, APIs are provided that make it easy to extend the library for your own use.

# The metaphor: Plugging components together into a board

# Main idea

Due to is modular, object-oriented design, gem5 can be thought of as a set of components that can be plugged together to form a simulation.

The main types of components are:

- **Component**: A component is a *model* with set parameters (or possibly a few customizable parameters). The types of components are *boards*, *processors*, *memory systems*, and *cache hierarchies*.
- **Board**: The "backbone" of the system. You plug components into the board. The board also contains the system-level things like devices, workload, etc. It's the boards job to negotiate the connections between other components.
- **Processor**: Processors connect to boards and have one or more *cores*
- **Cache hiearchy**: A cache hierarchy is a set of caches that can be connected to a processor and memory system.
- **Memory system**: A memory system is a set of memory controllers and memory devices that can be connected to the cache hierarchy.

# Quick note on relationship to gem5 models

The C++ code in gem5 specifies *parameterized* **models**.
These models are then instantiated in the Python scripts.

The standard library is a way to *wrap* these models in a standard API.
Most of the components in the standard library are models with pre-specified parameters.

If you want to change the values of the parameters of the models, you are encouraged to *extend* (i.e., subclass) the components in the standard library or create new components.
We will see some examples of this over the coming lectures.

# Let's get started!

In `materials/02-Using-gem5/01-stdlib/01-components.py` you'll see some imports already included for you.

```python
from gem5.components.boards.simple_board import SimpleBoard
from gem5.components.processors.simple_processor import SimpleProcessor
from gem5.components.cachehierarchies.ruby.mesi_two_level_cache_hierarchy import (
    MESITwoLevelCacheHierarchy,
)
from gem5.components.memory.single_channel import SingleChannelDDR4_2400
from gem5.components.processors.cpu_types import CPUTypes
from gem5.isas import ISA
from gem5.resources.resource import obtain_resource
from gem5.simulate.simulator import Simulator
```

# Let's build a system with a cache hierarchy

```
cache_hierarchy = MESITwoLevelCacheHierarchy(
    l1d_size="16kB",
    l1d_assoc=8,
    l1i_size="16kB",
    l1i_assoc=8,
    l2_size="256kB",
    l2_assoc=16,
    num_l2_banks=1,
)
```

`MESITwoLevelCacheHierarchy` is a component that represents a two-level MESI cache hierarchy.
This uses the [Ruby memory model](#).

The component for the cache hierarchy is parameterized with the sizes and associativities of the L1 and L2 caches.

# Next, let's add a memory system

```
memory = SingleChannelDDR3_1600()
```

This component represents a single-channel DDR3 memory system.

There is a `size` parameter that can be used to specify the size of the memory system of the simulated system. You can reduce the size to save simulation time, or use the default for the memory type (e.g., one channel of DDR3 defaults to 8 GiB).

There are also multi channel memories available.
We'll cover this more in Memory Systems.

# Next, let's add a processor

```
processor = SimpleProcessor(cpu_type=CPUTypes.TIMING, isa=ISA.ARM, num_cores=1)
```

`SimpleProcessor` is a component that allows you to customize the model for the underlying cores.

The `cpu_type` parameter specifies the type of CPU model to use.

# Next, plug components into the board

```
board = SimpleBoard(
    clk_freq="3GHz",
    processor=processor,
    memory=memory,
    cache_hierarchy=cache_hierarchy,
)
```

# Finally, set the workload and run the simulation

```
board.set_workload(obtain_resource("arm-gapbs-bfs-run"))

simulator = Simulator(board=board)
simulator.run()
```

## Run it

```
gem5-mesi 01-components.py
```

# Output

```
Generate Time:      0.00462
Build Time:         0.00142
Graph has 1024 nodes and 10496 undirected edges for degree: 10
Trial Time:         0.00010
Trial Time:         0.00008
Trial Time:         0.00008
Trial Time:         0.00008
Trial Time:         0.00008
Trial Time:         0.00009
Trial Time:         0.00008
Trial Time:         0.00008
Trial Time:         0.00008
Trial Time:         0.00011
Average Time:       0.00009
```

## stats.txt

```
simSeconds                              0.009093
simTicks                              9093461436
```

# Components included in gem5

```
gem5/src/python/gem5/components
----/boards
----/cachehierarchies
----/memory
----/processors

gem5/src/python/gem5/prebuilt
----/demo/x86_demo_board
----/riscvmatched
```

- gem5 stdlib in `src/python/gem5`
- Two types
  - Prebuilt: full systems with set parameters
  - Components: Components to build systems
- Prebuilt
  - Demo: Just examples to build off of
  - riscvmatched: Models SiFive Unmatched

# Components: Boards

```
gem5/src/python/gem5/components
----/boards
    ----/simple
    ----/arm_board
    ----/riscv_board
    ----/x86_board
----/cachehierarchies
----/memory
----/processors
```

- Boards: Things to plug into
  - Have "set_workload" and "connect_things"
  - Simple: SE-only, configurable
- Arm, RISC-V, and X86 versions for full system simulation

# Components: Cache hierarchies

```
gem5/src/python/gem5/components
----/boards
----/cachehierarchies
    ----/chi
    ----/classic
    ----/ruby
----/memory
----/processors
```

- Have fixed interface to processors and memory
- **Ruby:** detailed cache coherence and interconnect
- **CHI:** Arm CHI-based protocol implemented in Ruby
- **Classic caches:** Hierarchy of crossbars with inflexible coherence

# A bit more about cache hiearchies

- Quick caveat: You need different gem5 binaries for different protocols
- Any binary can use classic caches
- Only one Ruby protocol per gem5 binary

## In your codespaces, we have some pre-built binaries

- `gem5`: CHI (Fully configurable; based on Arm CHI)
- `gem5-mesi`: MESI_Two_Level (Private L1s, Shared L2)
- `gem5-vega`: GPU_VIPER (CPU: Private L1/L2 core pairs, shared L3; GPU: Private L1, shared L2)

# Components: Memory systems

```
gem5/src/python/gem5/components
----/boards
----/cachehierarchies
----/memory
     ----/single_channel
     ----/multi_channel
     ----/dramsim
     ----/dramsys
     ----/hbm
----/processors
```

- Preconfigured (LP)DDR3/4/5 DIMMs
  - Single and multi channel
- Integration with DRAMSim and DRAMSys
  - Not needed for accuracy, but useful for comparisons
- HBM: An HBM stack

# Components: Processors

```
gem5/src/python/gem5/components
----/boards
----/cachehierarchies
----/memory
----/processors
    ----/generators
    ----/simple
    ----/switchable
```

- Mostly "configurable" processors to build off
- Generators
  - Synthetic traffic, but act like processors
  - Have linear, random, and more interesting
- Simple
  - Only default parameters, one ISA
- Switchable
  - We'll see this later, but you can switch from one to another during simulation

# More on processors

- Processors are made up of cores
- Cores have a "BaseCPU" as a member
  - This is the actual CPU model
- `Processor` is what interfaces with `CacheHierarchy` and `Board`

**gem5 has three (or four or five) different processor models**

More details coming in the CPU Models section.

- `CPUTypes.TIMING`: A simple in-order CPU model
  - This is a "single cycle" CPU. Each instructions takes the time to fetch and executes immediately.
  - Memory operations take the latency of the memory system
  - OK for doing memory-centric studies, but not good for most research.

# CPU types

Other options for CPU types

- `CPUTypes.O3` : An out-of-order CPU model
  - Highly detailed model based on the Alpha 21264.
  - Has ROB, physical registers, LSQ, etc.
  - Don't use `SimpleProcessor` if you want to configure this.
- `CPUTypes.MINOR` : An in-order core model
  - A high-performance in-order core model.
  - Configurable four-stage pipeline
  - Don't use `SimpleProcessor` if you want to configure this.
- `CPUTypes.ATOMIC` : Used in "atomic" mode (more later)
- `CPUTypes.KVM` : More later

# Standard Library components

**Designed around *Extension* and *Encapsulation***
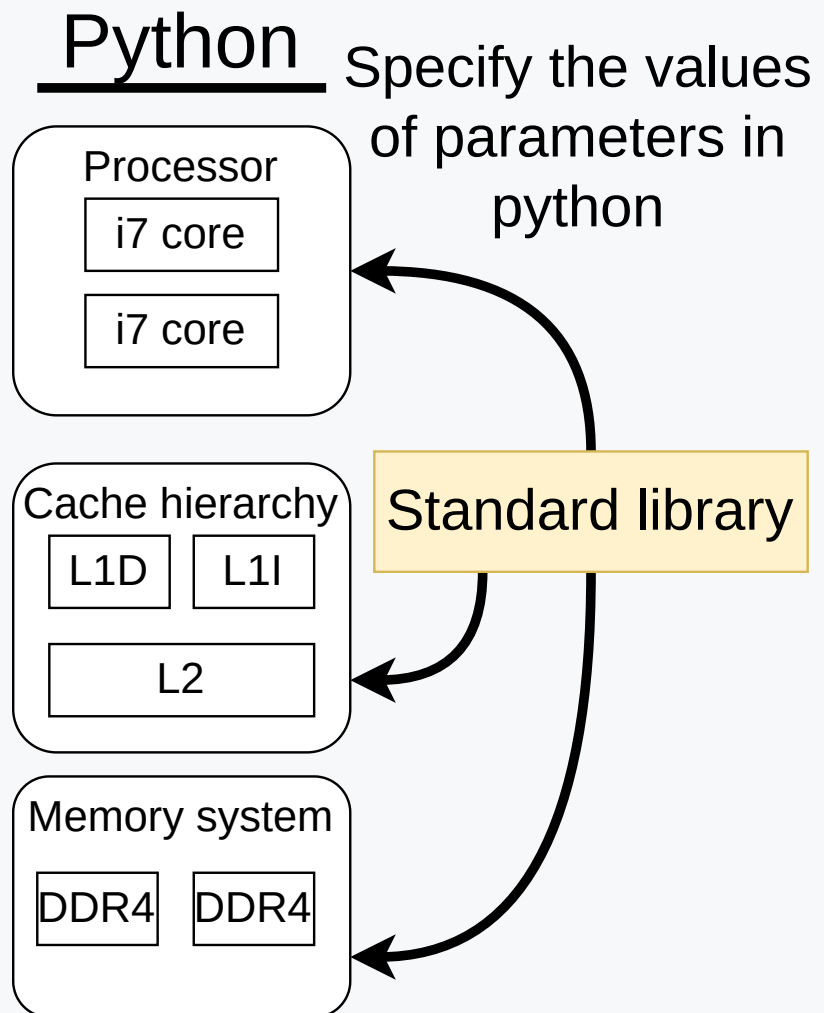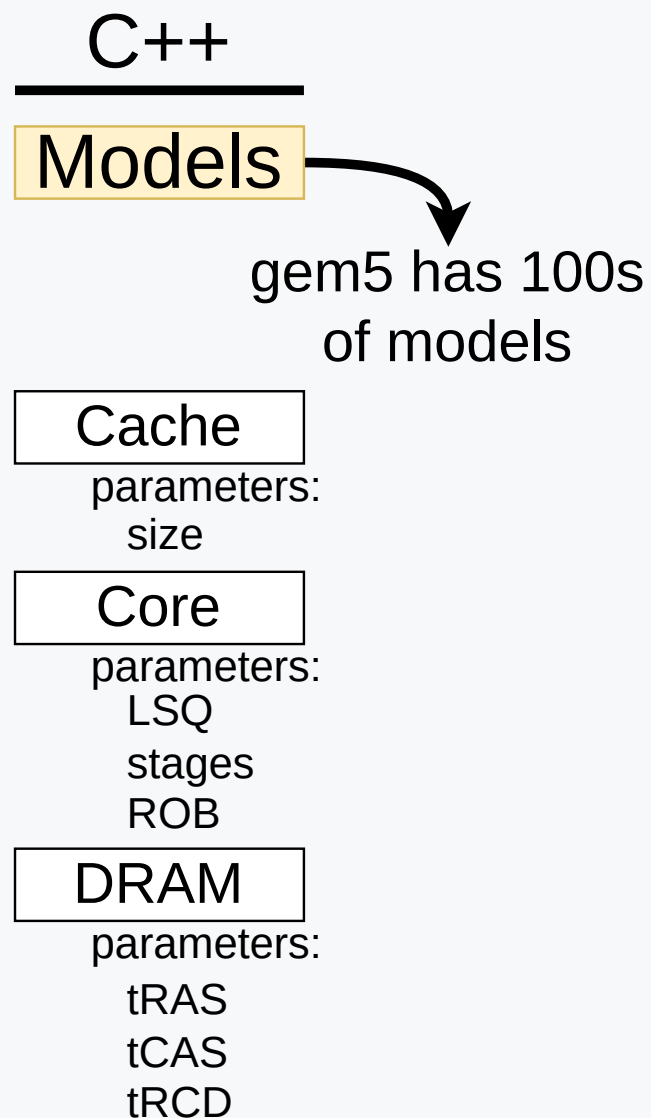
> NOT designed for "parameterization"

If you want to create a processor/cache hierarchy/etc. with different parameters:
Extend using object-oriented semantics

Let's see and example

# Quick reminder of gem5's architecture

We will now create a new component

Specialize/extend the "BaseO3CPU" (core)

## C++

**Models**

gem5 has 100s of models

### Cache
parameters:
size

### Core
parameters:
LSQ
stages
ROB

### DRAM
parameters:
tRAS
tCAS
tRCD

## Python

Specify the values of parameters in python

**Processor**
i7 core
i7 core

**Cache hierarchy**
L1D   L1I
L2

Standard library

**Memory system**
DDR4   DDR4

# Let's create a processor with out-of-order cores

Use `materials/02-Using-gem5/01-stdlib/02-processor.py` as a starting point.

Mostly the same as the last example, but now we have the following code instead of using the `SimpleProcessor`:

```
my_ooo_processor = MyOutOfOrderProcessor(
    width=8, rob_size=192, num_int_regs=256, num_fp_regs=256
)
```

# Create subclass of BaseCPUProcessor/Core

To specialize the parameters, create a subclass.

> Remember, don't just change the parameters of the `BaseCPUProcessor` directly.

```python
from m5.objects import ArmO3CPU
from m5.objects import TournamentBP
class MyOutOfOrderCore(BaseCPUCore):
    def __init__(self, width, rob_size,
                 num_int_regs, num_fp_regs):
        super().__init__(ArmO3CPU(), ISA.Arm)
```

See `...gem5/components/processors/base_cpu_core.py`
And `src/cpu/o3/BaseO3CPU.py`

```python
self.core.fetchWidth = width
self.core.decodeWidth = width
self.core.renameWidth = width
self.core.issueWidth = width
self.core.wbWidth = width
self.core.commitWidth = width

self.core.numROBEntries = rob_size
self.core.numPhysIntRegs = num_int_regs
self.core.numPhysFloatRegs = num_fp_regs

self.core.branchPred = TournamentBP()

self.core.LQEntries = 128
self.core.SQEntries = 128
```

# Now the `Processor`

The `CPUProcessor` assumes a list of cores that are `BaseCPUCores`

```python
class MyOutOfOrderProcessor(BaseCPUProcessor):
    def __init__(self, width, rob_size, num_int_regs, num_fp_regs):
        cores = [MyOutOfOrderCore(width, rob_size, num_int_regs, num_fp_regs)]
        super().__init__(cores)
```

We'll just make one core and we'll pass the parameters through to it.

# Now, run it and compare!

```
gem5-mesi 02-processor.py
```

Takes 2-3 minutes

## Questions

- Is this faster than simple in order?
- Use `--outdir=m5out/ooo` and `--outdir=simple`
- Compare the stats.txt (which stat?)

# Stat comparison

## Simple CPU

```
simSeconds                                              0.009073
board.processor.cores.core.ipc                          0.362353
```

## Our out-of-order CPU

```
simSeconds                                              0.003114
board.processor.cores.core.ipc                          1.055705
```

**Host seconds:** `17.09` vs `43.39`

O3 CPU takes more than 2x longer to simulate. More fidelity costs more time.

# Controlling simulation

# The `Simulator` object controls simulation

We'll see more about this in [Accelerating Simulation](#).

## Simulator parameters

- `board`: The `Board` to simulate (required)
- `full_system`: Whether to simulate a full system (default: `False`)
- `on_exit_event`: A complex data structure that allows you to control the simulation. The simulator exits for many reasons, this allows you to customize what happens. More on this later.
- `checkpoint_path`: If we're restoring from a checkpoint, this is the path to the checkpoint.
- `id`: An optional name for this simulation. Used in multisim and more in the future.

# Some useful functions

- `run()`: Run the simulation
- `get/set_max_ticks(max_tick)`: Set the absolute tick to stop simulation. Generates a `MAX_TICK` exit event that can be handled.
- `schedule_max_insts(inst_number)`: Set the number of instructions to run before stopping. Generates a `MAX_INSTS` exit event that can be handled. Note that if running multiple cores, this happens if *any* core reaches this number of instructions.
- **`get_stats()``**: Get the statistics from the simulation. Returns a dictionary of statistics.

See `src/python/gem5/simulate/simulator.py` for more details.

we will be covering much more about how to use the `Simulator` object in other parts of the bootcamp.

# Summary

- gem5's standard library is a set of components that can be used to build a simulation that does the majority of the work for you.
- The standard library is designed around *extension* and *encapsulation*.
- The main types of components are *boards*, *processors*, *memory systems*, and *cache hierarchies*.
- The standard library is designed to be modular and object-oriented.
- The `Simulator` object controls the simulation.