Modeling CPU cores in gem5

Outline

- **Learn about CPU models in gem5**
  - AtomicSimpleCPU, TimingSimpleCPU, O3CPU, MinorCPU, KvmCPU
- Using the CPU models
  - Set-up a simple system with two cache sizes and three CPU models
- Look at the gem5 generated statistics
  - To understand differences among CPU models
- Create a custom processor
  - Change parameters of a processor based on O3CPU

gem5

gem5 CPU Models

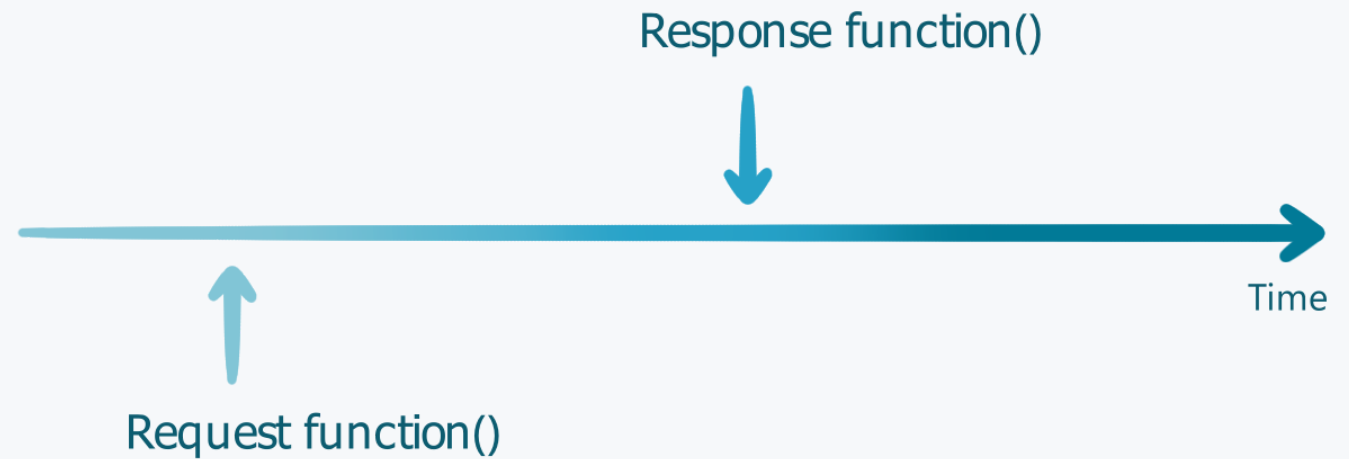# Simple CPU

SimpleCPU

Atomic

Sequency of nested calls
Use: Warming up, fast-forwarding

Functional

Backdoor access to mem.
(loading binaries)
No effect on coherency states

Timing

Split transactions
Models queuing delay and
resource contention

Return of Request function() with an *approximate time*

Time

Request function()

Response function()

Time

Request function()

Other Simple CPUs

AtomicSimpleCPU

- Uses *Atomic* memory accesses
  - No resource contentions or queuing delay
  - Mostly used for fast-forwarding and warming of caches

TimingSimpleCPU

- Uses *Timing* memory accesses
  - Execute non-memory operations in one cycle
  - Models the timing of memory accesses in detail

O3CPU (Out of Order CPU Model)

- *Timing* memory accesses *execute-in-execute* semantics
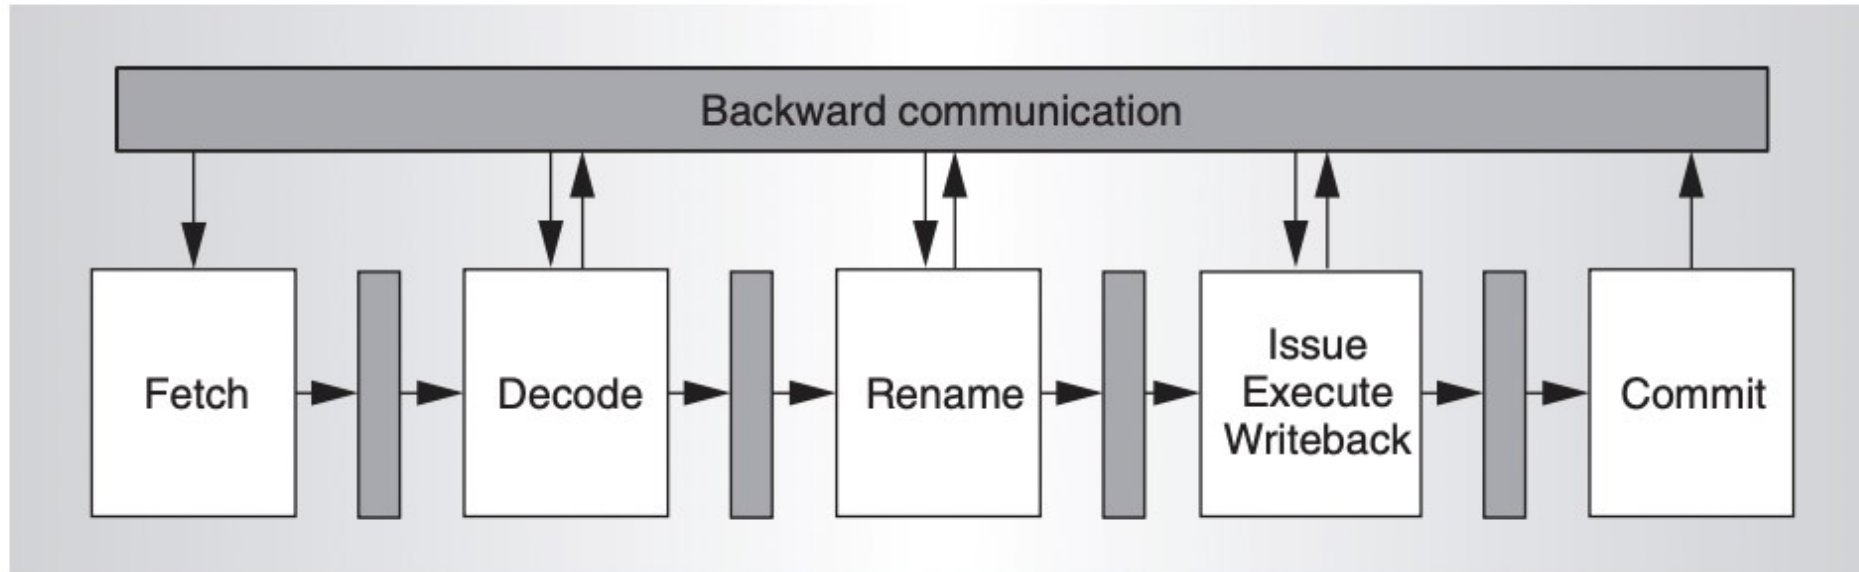- Time buffers between stages

Figure 2. O3CPU pipeline. Shaded boxes represent time buffers.

The O3CPU Model has many parameters

src/cpu/o3/BaseO3CPU.py

```
decodeToFetchDelay = Param.Cycles(1, "Decode to fetch delay")
renameToFetchDelay = Param.Cycles(1, "Rename to fetch delay")
...
fetchWidth = Param.Unsigned(8, "Fetch width")
fetchBufferSize = Param.Unsigned(64, "Fetch buffer size in bytes")
fetchQueueSize = Param.Unsigned(
    32, "Fetch queue size in micro-ops per-thread"
)
...
```

Remember, do not update the parameters directly in the file. Instead, create a new *stdlib component* and extend the model with new values for parameters.
We will do this soon.

MinorCPU

KvmCPU

- KVM – Kernel-based virtual machine
- Used for native execution on x86 and ARM host platforms
- Guest and the host need to have the same ISA
- Very useful for functional tests and fast-forwarding

gem5

Summary of gem5 CPU Models

BaseKvmCPU

- Very fast
- No timing
- No caches, BP

BaseSimpleCPU

- Fast
- Some timing
- Caches, limited BP

DerivO3CPU and MinorCPU

- Slow
- Timing
- Caches, BP

BaseCPU

BaseKvmCPU

ArmV8Kvm

X86Kvm

BaseSimpleCPU

TimingSimple

AtomicSimple

DerivO3CPU

MinorCPU

gem5

Interaction of CPU model with other parts of gem5

Outline

- Learn about CPU models in gem5
  - AtomicSimpleCPU, TimingSimpleCPU, O3CPU, MinorCPU, KvmCPU
- **Using the CPU models**
  - Set-up a simple system with two cache sizes and three CPU models
- Look at the gem5 generated statistics
  - To understand differences among CPU models
- Create a custom processor
  - Change parameters of a processor based on O3CPU

Let's use these CPU Models!

gem5

Material to use

Start by opening the following file.

materials/02-Using-gem5/04-cores/cores.py

Steps

1. Configure a simple system with Atomic CPU
2. Configure the same system with Timing CPU
3. Reduce the cache size
4. Change the CPU type back to Atomic

We will be running a workload called matrix-multiply on **different CPU types and cache sizes**.

Let's configure a simple system with Atomic CPU

materials/02-Using-gem5/04-cores/cores.py

```python
from gem5.components.boards.simple_board import SimpleBoard
from gem5.components.cachehierarchies.classic.private_l1_cache_hierarchy import PrivateL1CacheHierarchy
from gem5.components.memory.single_channel import SingleChannelDDR3_1600
from gem5.components.processors.cpu_types import CPUTypes
from gem5.components.processors.simple_processor import SimpleProcessor
from gem5.isas import ISA
from gem5.resources.resource import obtain_resource
from gem5.simulate.simulator import Simulator


# A simple script to test with different CPU models
# We will run a simple application (matrix-multiply) with AtomicSimpleCPU and TimingSimpleCPU
# using two different cache sizes
...
```

Let's start with Atomic CPU

`cpu_type` in cores.py should already be set to Atomic.

```
# By default, use Atomic CPU
cpu_type = CPUTypes.ATOMIC


# Uncomment for steps 2 and 3
# cpu_type = CPUTypes.TIMING
```

Let's run it!

```
gem5 --outdir=atomic-normal-cache ./materials/02-Using-gem5/04-cores/cores.py
```

Make sure the out directory is set to **atomic-normal-cache**.

Next, try Timing CPU

Change `cpu_type` in cores.py to Timing.

```
# By default, use Atomic CPU
# cpu_type = CPUTypes.ATOMIC


# Uncomment for steps 2 and 3
cpu_type = CPUTypes.TIMING
```

Let's run it!

```
gem5 --outdir=timing-normal-cache ./materials/02-Using-gem5/04-cores/cores.py
```

Make sure the out directory is set to **timing-normal-cache**.

Now, try changing the Cache Size

Go to this line of code.

```
cache_hierarchy = PrivateL1CacheHierarchy(l1d_size="32KiB", l1i_size="32KiB")
```

Change `l1d_size` and `l1i_size` to 1KiB.

```
cache_hierarchy = PrivateL1CacheHierarchy(l1d_size="1KiB", l1i_size="1KiB")
```

Let's run it!

```
gem5 --outdir=timing-small-cache ./materials/02-Using-gem5/04-cores/cores.py
```

Make sure the out directory is set to **timing-small-cache**.

Now let's try a Small Cache with Atomic CPU

Set `cpu_type` in cores.py to Atomic.

```
# By default, use Atomic CPU
cpu_type = CPUTypes.ATOMIC


# Uncomment for steps 2 and 3
# cpu_type = CPUTypes.TIMING
```

Let's run it!

```
gem5 --outdir=atomic-small-cache ./materials/02-Using-gem5/04-cores/cores.py
```

Make sure the out directory is set to **atomic-small-cache**.

Outline

- Learn about CPU models in gem5
  - AtomicSimpleCPU, TimingSimpleCPU, O3CPU, MinorCPU, KvmCPU

- Using the CPU models
  - Set-up a simple system with two cache sizes and three CPU models

- **Look at the gem5 generated statistics**
  - To understand differences among CPU models

- Create a custom processor
  - Change parameters of a processor based on O3CPU

Statistics

Look at the Number of Operations

Run the following command.

```
grep -ri "simOps" *cache
```

Here are the expected results. (Note: Some text is removed for readability.)

```
atomic-normal-cache/stats.txt:simOps                    33954560
atomic-small-cache/stats.txt:simOps                     33954560
timing-normal-cache/stats.txt:simOps                    33954560
timing-small-cache/stats.txt:simOps                     33954560
```

"Ops" may be different from "Instructions" because gem5 breaks instructions down into "micro-ops."
x86 is highly microcoded, all ISAs have some microcoded instructions in gem5.

Look at the Number of Execution Cycles

Run the following command.

```
grep -ri "cores0.*numCycles" *cache
```

Here are the expected results. (Note: Some text is removed for readability.)

```
atomic-normal-cache/stats.txt:board.processor.cores0.core.numCycles        38157549
atomic-small-cache/stats.txt:board.processor.cores0.core.numCycles         38157549
timing-normal-cache/stats.txt:board.processor.cores0.core.numCycles        62838389
timing-small-cache/stats.txt:board.processor.cores0.core.numCycles         96494522
```

Note that for Atomic CPU, the number of cycles is the **same** for a large cache *and* a small cache.

This is because Atomic CPU ignores memory access latency.

Extra Notes about gem5 Statistics

When you specify the out-directory for the stats file (when you use the flag `--outdir=<outdir-name>`), go to **<outdir-name>/stats.txt** to look at the entire statistics file.

For example, to look at the statistics file for the Atomic CPU with a small cache, go to **atomic-small-cache/stats.txt**.

In general, if you don't specify the out-directory, it will be **m5out/stats.txt**.

Other statistics to look at

- Host time (time taken by gem5 to run your simulation)
  - *hostSeconds*

Outline

- Learn about CPU models in gem5
  - AtomicSimpleCPU, TimingSimpleCPU, O3CPU, MinorCPU, KvmCPU
- Using the CPU models
  - Set-up a simple system with two cache sizes and three CPU models
- Look at the gem5 generated statistics
  - To understand differences among CPU models
- **Create a custom processor**
  - Change parameters of a processor based on O3CPU

Let's configure a custom processor!

Material to use

materials/02-Using-gem5/04-cores/cores-complex.py

materials/02-Using-gem5/04-cores/components/processors.py

Steps

1. Update class Big(O3CPU) and Little(O3CPU)
2. Run with Big processor
3. Run with Little processor
4. Compare statistics

We will be running the same workload (matrix-multiply) on **two custom processors**.

Configuring two processors

We will make one fast processor (**Big**) and one slow processor (**Little**).

To do this, we will change **4** parameters in each processor.

- **width**
  - width of fetch, decode, rename, issue, wb, and commit stages
- **rob_size**
  - the number of entries in the reorder buffer
- **num_int_regs**
  - the number of physical integer registers
- **num_fp_regs**
  - the number of physical vector/floating point registers

Configuring Big

Open the following file.
materials/02-Using-gem5/04-cores/components/processors.py

In class Big, set

- width=**10**
- rob_size=**40**
- num_int_regs=**50**
- num_fp_regs=**50**

```python
class Big(O3CPU):
    def __init__(self):
        super().__init__(
            width=0,
            rob_size=0,
            num_int_regs=0,
            num_fp_regs=0,
        )
```

Configuring Little

Keep working in the following file.
materials/02-Using-gem5/04-cores/components/processors.py

In class Little, set

- width=**2**
- rob_size=**30**
- num_int_regs=**40**
- num_fp_regs=**40**

```python
class Little(O3CPU):
    def __init__(self):
        super().__init__(
            width=0,
            rob_size=0,
            num_int_regs=0,
            num_fp_regs=0,
        )
```

Run with Big processor

We will be running the following file.
materials/02-Using-gem5/04-cores/cores-complex.py

First, we will run matrix-multiply with our Big processor.

Run with the following command.

```
gem5 --outdir=big-proc ./materials/02-Using-gem5/04-cores/cores-complex.py -p big
```

Make sure the out directory is set to **big-proc**.

Run with Little processor

Next, we will run matrix-multiply with our Little processor.

Run with the following command.

```
gem5 --outdir=little-proc ./materials/02-Using-gem5/04-cores/cores-complex.py -p little
```

Make sure the out directory is set to **little-proc**.

Comparing Big and Little processors

Run the following command.

```
grep -ri "simSeconds" *proc && grep -ri "numCycles" *proc
```

Here are the expected results. (Note: Some text is removed for readability.)

```
big-proc/stats.txt:simSeconds                                    0.028124
little-proc/stats.txt:simSeconds                                 0.036715
big-proc/stats.txt:board.processor.cores.core.numCycles          56247195
little-proc/stats.txt:board.processor.cores.core.numCycles       73430220
```

Our Little processor takes more time and more cycles than our Big processor.