

**Sampled
simulation with
gem5**



Recap: What if the ROI is large

We now know how to skip the "unimportant" part of the simulation, but what if the important part of the simulation is too too large?

Sampling

What is sampling

There are two major types of sampling:

1. Targeted sampling
2. Statistical sampling

Targeted Sampling

Representative methodologies: SimPoint, LoopPoint




Targeted sampling selects samples based on specific characteristics that are discovered by analysis.

More about Targeted Sampling

Well-known simulation sampling methods that use the targeted sampling approach include **SimPoint** and **LoopPoint**.

Both methods divided the whole program execution into regions that each executes a fixed number of instructions.


They record the basic block execution pattern inside each region to be the signature of the program behavior for the region. It's known as the **basic block vector**. Here is an example of a basic block vector.



They use the basic block vectors to cluster and find the representative regions. They will predict the overall performance of the program by collecting only the representative regions' performance and sum them with weights.

Statistical Sampling

Representative methodologies: SMARTS, FSA



Statistical sampling, as the name suggests, statistically selects its sampling units.

More about Statistical Sampling

The representative simulation sampling methods in statistical sampling include **SMARTS** and **FSA**.

Both methods periodically or randomly simulated in detailed for a small amount of execution throughout the whole program execution, and fast-forward between the detailed simulations. They use the performance of the randomly distributed samples to predict the overall performance of the whole program execution.

What should we know before we apply the techniques

No matter how great a tool or a technique is, misusing it can be **DANGEROUS**

Before using any of the sampling techniques, we need to make sure the sampling technique works for our experiments.

For example, SimPoint is designed to work with single-threaded workloads only, so **if our experiments require multi-threaded workloads, we should NOT use SimPoint with them.**



What gem5 offers

In gem5, we have infrastructures for

1. SimPoint ([paper](#))
2. LoopPoint ([paper](#))
3. ElFie ([paper](#))
4. SMARTS ([paper](#))
5. FSA ([paper](#)) (Might not be supported officially)



Targeted Sampling in gem5



Targeted Sampling in gem5

- gem5 provides infrastructures for **SimPoint** and **LoopPoint** to analyze the program, take checkpoints for the representative regions, and run the representative regions.
Note that LoopPoint's analysis support is not currently supported in gem5 24.0 but is tested and prepared to be upstream to gem5 24.1.
- gem5 also provides infrastructures for **ElFies** to be executed in SE mode, but gem5 does not support creating ElFies files and information.



SimPoint in gem5



SimPoint

As mentioned before, there are three steps in using SimPoint:

1. Analysis
2. Taking checkpoints
3. Run the regions

There are two key files related to using SimPoint in gem5:

1. [src/python/gem5/utils/simpoint.py](#)
2. [src/cpu/simple/probes/SimPoint.py](#)

We will step through them in this section

SimPoint Analysis

In gem5, we use the `SimPoint` probe listener object to collect the information SimPoint needed to cluster the regions.

This object is defined in <src/cpu/simple/probes/SimPoint.py>.

The `SimPoint` probe listener has two parameters: [interval](#) and [profile_file](#).

- The `interval` takes in a length as our definition of a region. It means that every time we execute n number of instructions, we see it as the end of a region. The default length is 100,000,000.
- The `profile_file` takes in a name for the output zip file. The default name is "simpoint.bb.gz".

In order to use this probe listener object, we need to attach it to an ATOMIC CPU. It will start collecting information as soon as the simulation starts and stops when the simulation ends. After exiting the simulation, there will be a zip file with the basic block vector information for each region under the simulation outdir directory.

Hands-On Time!

01-simpoint

All materials can be found under `materials/02-Using-gem5/09-sampling/01-simpoint`. The completed version is under `materials/02-Using-gem5/09-sampling/01-simpoint/complete`.

In this exercise, we will not modify any of the scripts, but just to run the scripts.

We will only work with [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint-analysis.py](#) and [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint3.2-cmd.sh](#).

Goal

1. Run the SimPoint analysis
2. Process the data to get the representative regions and their weights



01-simpoint

Because the profiling might take a while, so we can run the simulation first with the following command

```
gem5 -re --outdir=full-detailed-run-m5out full-detailed-run.py  
gem5 -re --outdir=simpoint-analysis-m5out simpoint-analysis.py
```

In this exercise, we are trying to create SimPoints for a simple workload.

The source code of the simple workload can be found in [materials/02-Using-gem5/09-sampling/01-simpoint/workload/simple_workload.c](#).

It allocates an array of one thousand 64 bits elements, assign each a number, then sum all up with one thousand iterations.

We can expect the program behavior of this workload to be really repeating.

01-simpoint

The [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint-analysis.py](#) uses the `SimPoint` probe listener object that we introduced earlier to collect basic block information for this simple workload.

It connects the ATOMIC CPU core to the `SimPoint` probe listener using

```
processor.get_cores()[0].core.addSimPointProbe(1_000_000)
```

The definition of this `addSimPointProbe()` function can be found under [src/cpu/simple/BaseAtomicSimpleCPU.py](#).

In this example, we set the `interval_length` as 1,000,000, which means that we define a region as 1,000,000 instruction executed (committed).

01-simpoint

After the simulation finishes, we will see a zip file named `simpoint.bb.gz` under the `simpoint-analysis-m5out` folder.

We can unzip it with the following command

```
gzip -d -k simpoint.bb.gz
```

After unzipping it, we can look at the `simpoint.bb` file.

This file contains all basic block vector information for the simple workload.

There are 9 regions found in this workload.

Each region has a basic block vector, which starts with a `T`.

01-simpoint

As we can see, region 2 to region 9 have almost identical basic block vectors.

```
T:1900:222 :1901:222 :1902:999216 :1903:333
T:1900:222 :1901:222 :1902:999225 :1903:333
T:1900:222 :1901:222 :1902:999225 :1903:333
T:1900:222 :1901:222 :1902:999225 :1903:333
T:1900:222 :1901:222 :1902:999216 :1903:333
T:1900:222 :1901:222 :1902:999225 :1903:333
T:1900:222 :1901:222 :1902:999225 :1903:333
T:1900:222 :1901:222 :1902:999225 :1903:333
```

The similar basic block vectors indicate the program behavior for region 2 to region 9 are really similar. Therefore, we can expect to cluster region 2 to region 9 together, and pick one region out of it to be our representative region to represent the performance of all regions from region 2 to region 9.

01-simpoint

Let's further understand what this line means

```
T:1900:222 :1901:222 :1902:999216 :1903:333
```

- **T** means the start of the region's basic block vector.
- **:1900:222** means basic block 1900 executed (committed) 222 instructions. The key here is that 222 is NOT the number of time the basic block has been executed, but the number of time the basic block has been executed multiples with the basic block's total instructions. If we sum up all the instructions that have been executed by the basic blocks, we will get roughly the length of the region. $222+222+999216+333=999993$.

The next step is to use this information to cluster the regions and find the representative regions. There are many methods to do it. In this exercise, we will be using the SimPoint3.2 tool that was provided by the SimPoint paper authors.

01-simpoint

The tool is already compiled under [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint](#). We also provided a runscript with the command in [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint3.2-cmd.sh](#).

```
/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/simpoint \  
-inputVectorsGzipped -loadFVFile simpoint-analysis-m5out/simpoint.bb.gz -k 5 -saveSimpoints \  
results.simpts -saveSimpointWeights results.weights
```

Let's look at this command.

It pass in the `simpoint.bb.gz` to the tool.

It set the number of clusters expected to be 5 using `-k 5`.

It saves the SimPoint information in `results.simpts` and their weights in `results.weight`.

01-simpoint

After running command, we will get the following file

```
#results.simpts`           # results.weights
2 0                        0.666667 0
1 1                        0.222222 1
0 3                        0.111111 3
```

This means that it found 3 SimPoints out of this program.
Region 2 is SimPoint 0, which has a weight of 0.666667.
Region 1 is SimPoint 1, which has a weight of 0.222222.
Region 0 is SimPoint 3, which has a weight of 0.111111.

There are not 5 clusters because the algorithm found that 3 clusters are enough to represent all the program behaviors.

The SimPoint tag number might not be continuous because it is the tag number for the cluster.

SimPoint Checkpoint

Now we have the representative regions and their weights, we will need to find a way to get to those SimPoints.

As introduced in [08-accelerating-simulation](#), there are two methods to get to the region of interest. With SimPoint, we usually use checkpointing.

In order to do it easily in gem5, we will be using the [set_se_simpoint_worload](#), the [SimPoint class](#), and the [simpoints_save_checkpoint_generator](#).

Hands-On Time!

01-simpoint

All materials can be found under `materials/02-Using-gem5/09-sampling/01-simpoint`. The completed version is under `materials/02-Using-gem5/09-sampling/01-simpoint/complete`.

We will still not modify and scripts.

Goal

- Use the output file from the SimPoint3.2 tool to create checkpoints for all 3 SimPoints

Because it might take some time to take the checkpoints, let's start taking checkpoints with the following commands

```
gem5 -re --outdir=simpoint-checkpoint-m5out simpoint-checkpiont.py
```

01-simpoint

Let's look at [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint-checkpoint.py](#).

There are three key parts

```
# key 1: pass in the representative regions information
from gem5.utils.simpoint import SimPoint
simpoint_info = SimPoint(
    simpoint_interval=1_000_000,
    simpoint_file_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/results.simpts"),
    weight_file_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/results.weights"),
    warmup_interval=1_000_000
)
# key 2: pass in the SimPoint object to setup the workload
board.set_se_simpoint_workload(
    binary=BinaryResource(local_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/workload/simple_workload").as_posix()),
    simpoint=simpoint_info
)
# key 3: register exit event handler to take the checkpoints
simulator = Simulator(
    board=board,
    on_exit_event={
        ExitEvent.SIMPOINT_BEGIN: simpoints_save_checkpoint_generator(dir, simpoint_info)
    },
)
```

01-simpoint

Let's zoom in to key 1.

```
from gem5.utils.simpoint import SimPoint
simpoint_info = SimPoint(
    simpoint_interval=1_000_000,
    simpoint_file_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/results.simpts"),
    weight_file_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/results.weights"),
    warmup_interval=1_000_000
)
```

The name of the parameters for `simpoint_interval`, `simpoint_file_path`, and `weight_file_path` should speak for themselves, but `warmup_interval` is new to us in this section.

As mentioned in [08-accelerating-simulation](#), we should expect most of the micro-architectural states to be cold when restoring a checkpoint, therefore, we need to reserve a warm up period for our region of interest to warm up the micro-architectural states of the system.

Therefore, we need to setup a `warmup_interval` value here, so when we restore the checkpoint, we can expect to have this length of simulation to warm up our micro-architectural states.



01-simpoint

The SimPoint class can also detect how much space we have for the warm up period, if there is not enough space, it will automatically shrink the warm up period of that SimPoint to the space available. For example, we have a SimPoint that starts at the beginning of the program, then the warm up period for it should be 0 instruction.

The SimPoint class also automatically sort the SimPoints based on when it starts in terms of instructions.

It also provides getter functions that might be helpful, such as `get_weight_list()` that returns the list of the weights for each SimPoint.

More can be found in the [SimPoint](#) class definition.

01-simpoint

For key 2 and key 3, let's just trust them, because they are all relied on the information passed in from the `SimPoint` object.

We can declare where we want to store the SimPoint checkpoints by passing in the destination path to the `simpoints_save_checkpoint_generator` as a parameter. For more details, we can find the generator in src/python/gem5/simulate/exit_event_generators.py.

01-simpoint

After finish checkpointing, we should find all the SimPoint checkpoints under the `materials/02-Using-gem5/09-sampling/01-simpoint/simpoint-checkpoint` because we pass in `simpoint-checkpoint` to the `simpoints_save_checkpoint_generator` as our directory to save the SimPoint checkpoints.

There should be three checkpoint folders called `cpt.SimPoint0`, `cpt.SimPoint1`, and `cpt.SimPoint2`.

We can now use them to run our SimPoints for the simple workload.

Running SimPoint

Now we have the representative SimPoint checkpoint, we can run them with the systems we want to measure and predict the overall performance of running the whole simple workload.

SimPoint relies on the weights we got from the analysis stage to do the prediction.
The weight is calculated with

$$\text{Weight of cluster } i = \frac{\text{Number of regions in cluster } i}{\text{Total number of regions in the workload execution}}$$

For example, if we want to calculate the predicted IPC, we should

$$\text{Predicted IPC} = \sum_{i=1}^n (\text{Weight of cluster } i \times \text{IPC of cluster } i)$$

01-simpoint

All materials can be found under `materials/02-Using-gem5/09-sampling/01-simpoint`. The completed version is under `materials/02-Using-gem5/09-sampling/01-simpoint/complete`.

We will still not modify and scripts.

Unlike the simple systems we used for SimPoint analysis and SimPoint checkpointing, we now need to use the systems we actual want to measure the performance of.

We will be using [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint-run.py](#) to run our SimPoints.

For our baseline, we are using [materials/02-Using-gem5/09-sampling/01-simpoint/full-detailed-run.py](#), which run the whole simple workload with the detailed system.

Let's start running the SimPoints first before explaining how it works due to time constraint.

We provided a runsript to run all three in [materials/02-Using-gem5/09-sampling/01-simpoint/run-all-simpoint.sh](#)

```
./run-all-simpoint.sh
```


01-simpoint

Let's look at [materials/02-Using-gem5/09-sampling/01-simpoint/simpoint-run.py](#). It has a detailed system that matches the one that is used in our baseline [materials/02-Using-gem5/09-sampling/01-simpoint/full-detailed-run.py](#).

There are a few keys that we want to look at. Let's start with the one we are familiar with

```
# key 1:
from gem5.utils.simpoint import SimPoint
simpoint_info = SimPoint(
    simpoint_interval=1_000_000,
    simpoint_file_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/results.simpts"),
    weight_file_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/results.weights"),
    warmup_interval=1_000_000
)
# key 2:
board.set_se_simpoint_workload(
    binary=BinaryResource(local_path=Path("/workspaces/2024/materials/02-Using-gem5/09-sampling/01-simpoint/workload/simple_workload").as_posix()),
    simpoint=simpoint_info,
    checkpoint=Path(f"simpoint-checkpoint/cpt.SimPoint{args.sid}")
)
```

In here, we pass in the SimPoint information and use them to setup a workload that restore the checkpoint with our target SimPoint id.



01-simpoint

```
# key 3:
def max_inst():
    warmed_up = False
    while True:
        if warmed_up:
            print("end of SimPoint interval")
            yield True
        else:
            print("end of warmup, starting to simulate SimPoint")
            warmed_up = True
            # Schedule a MAX_INSTS exit event during the simulation
            simulator.schedule_max_insts(
                board.get_simpoint().get_simpoint_interval()
            )
            m5.stats.dump()
            m5.stats.reset()
            yield False

simulator = Simulator(
    board=board,
    on_exit_event={ExitEvent.MAX_INSTS: max_inst()},
)
```

Here is an exit event handler we defined for handling warm up period and detailed simulation period. It schedules an end for the SimPoint after the warm up period, and dump and reset the stats for detailed measurement.

01-simpoint

```
# key 4:
warmup_interval = board.get_simpoint().get_warmup_list()[args.sid]
if warmup_interval == 0:
    warmup_interval = 1
print(f"Starting Simulation with warmup interval {warmup_interval}")
simulator.schedule_max_insts(warmup_interval)
simulator.run()

print("Simulation Done")
print(f"Ran SimPoint {args.sid} with weight {board.get_simpoint().get_weight_list()[args.sid]}")
```

Before the starting the simulation, we need to setup an exit event to indicates the end of the warm up period.

We use the `simulator.schedule_max_insts()` function to do so.

We can use the `get_warmup_list()` from the `SimPoint` object to get the warm up period length for each `SimPoint`.

Here is one limitation of the `simulator.schedule_max_insts()` function, if it gets a value `0`, there won't be any exit event scheduled, so if the warm up period length is `0`, we have to set it to `1`.

01-simpoint

After setting up the above keys, the script is ready to run the SimPoint. Note that each simulation can only run one SimPoint.

After the running the SimPoints, we should see the output folder `simpoint0-run`, `simpoint1-run`, and `simpoint2-run`.

Also, we should have the baseline output in `full-detailed-run-m5out`.

Let's try to use the SimPoints' performance to predict the overall IPC!

We have a python script ready to do the prediction in [materials/02-Using-gem5/09-sampling/01-simpoint/predict_overall_ipc.py](#).

We can run it with

```
python3 predict_overall_ipc.py
```

01-simpoint

We should see something like this

```
predicted IPC: 1.2577933618669999  
actual IPC: 1.247741  
relative error: 0.8056449108428648%
```

The python script reads the IPC from our baseline.
It also reads the detailed simulation period IPC from all our SimPoints' stats files.
Then it performs the calculation with

```
predicted_ipc = 0.0  
for i in range(num_simpoints):  
    predicted_ipc += simpoint_ipcs[i] * simpoint_weights[i]
```

As the output suggests, the relative error between the predicted IPC and the actual baseline IPC is around 0.81%.

Summary in SimPoint

Congratulation! We walked through the whole process of how to do sampling with the SimPoint method.

Let's recap what we did:

1. Analyze the program
2. Checkpointing the representative regions
3. Run the regions and predict the performance

Here is a good news, this process is really similar for majority of the methods in targeted sampling. Therefore, if we know how to do sampling with the SimPoint method, it should not be hard for us to do others, such as LoopPoint, which enables multi-threaded workload sampling.

LoopPoint and ElFies in gem5

LoopPoint and Elfies

LoopPoint is similar to SimPoint with some key differences.

LoopPoint uses the number of time a loop is executed to mark the regions instead of using instruction executed. Therefore, we need collect the loop execution information in the analysis stage beside the basic block execution information.

Other than this, in terms of the process, it is really similar to SimPoint.

We will not do a detailed example on LoopPoint, but there are example under the [configs/example/gem5_library/looppoints](#). We have the infrastructures ready for the LoopPoint checkpointing and running the LoopPoint in gem5 v24.0. The LoopPoint analysis is tested and ready to be upstream in gem5 v24.1.

Elfies is a checkpointing method that creates checkpoint executables out of a large workload execution. It can be used with LoopPoint to create executables of the representative regions.

Like LoopPoint, it uses the number of time a loop has been executed to mark the beginning and the end of the region of interest, so we need those information to execute Elfies in gem5.

[Link](#) for more information about LoopPoint and Elfies



ElFies

gem5 does not produce ElFies, but we have support to run ElFies in SE mode. For all the weights and loop information, they should come with the ElFies workload.

We can run ElFies with the [ElFieInfo class](#).

There is an example in [materials/02-Using-gem5/09-sampling/02-elfies/run-elfies.py](#)

We can run it with the following command but it is not suggested because it will take too long.

```
gem5 -re run-elfies.py
```

It is a 8-threaded experiments with a detailed system that might run for eight hundred million instructions so it will take sometime to finish.

We have a completed m5out under the `materials/02-Using-gem5/09-sampling/02-elfies/complete/m5out` if you are interested in looking at the output.



ElFies Example

In the gem5 resource, we have some ElFies workload provided, for example, [wrf-s.1_globalr13](#). Its marker (loop execution) information can be found in [elfie-info-wrf-s.1_globalr13](#).

We can setup the ElFies workload as a normal executable

```
board.set_se_binary_workload(  
    binary=obtain_resource("wrf-s.1_globalr13")  
)
```

Then, we can setup the beginning and end marker with

```
from gem5.resources.elfie import ELFieInfo  
elfie = ELFieInfo(start = PcCountPair(int("0x100b643",0),1), end = PcCountPair(int("0x526730",0),297879) )  
elfie.setup_processor(  
    processor = processor  
)
```

ElFies Example

After getting to the beginning and end markers, the simulator will raise an exit event `SIMPOINT_BEGIN`, but there is not a default exit event handler for it, so we will need to define ours.

```
def start_end_handler():
    # This is a generator to handle exit event produced by the
    # start marker and end marker.
    # When we reach the start marker, we reset the stats and
    # continue the simulation.
    print(f"reached {targets[0]}\n")
    print("now reset stats\n")
    m5.stats.reset()
    print("fall back to simulation\n")
    yield False
    # When we reach the end marker, we dump the stats to the stats file
    # and exit the simulation.
    print(f"reached {targets[1]}\n")
    print("now dump stats and exit simulation\n")
    m5.stats.dump()
    yield True

simulator = Simulator(
    board=board,
    on_exit_event={
        ExitEvent.SIMPOINT_BEGIN : start_end_handler()
    }
)
```

Summary in LoopPoint and Elfies

After finish running all the Elfies for a macro-benchmark, we can use the weights the Elfies files provided to predict the overall performance like we did for the SimPoint example.

Now we covered all the targeted sampling methods that are supported in gem5, let's dive in to Statistical Sampling!



Statistical Sampling in gem5



SMARTS

SMARTS is one of the statical sampling methods.

It uses the statical model to predict the overall performance with randomly or periodically selected samples.

We use the small samples distributed throughout the whole program execution to predict the average performance. We can also use the average performance to predict the overall performance, such as runtime.

Before running the simulation, we need to decide a few statistical parameters

- n : the number of the samples. It is a number of count.
- k : the systematic sampling interval. It is a number of count.
- U : the sampling unit size. It is a number of the instruction executed.
- W : the length of the detailed warm up period. It is a number of the instruction executed.

Hands-On Time

03-SMARTS

All materials can be found under `materials/02-Using-gem5/09-sampling/03-SMARTS`. The completed version is under `materials/02-Using-gem5/09-sampling/03-SMARTS/complete`.

We will still not modify and scripts.

[materials/02-Using-gem5/09-sampling/03-SMARTS/SMARTS.py](#) is an example of how to use exit event handler to perform SMARTS.

We can run it with the following command

```
gem5 -re SMARTS.py
```

This script will use SMARTS on the workload from 01-simpiont, so we can use the baseline performance from the [materials/02-Using-gem5/09-sampling/01-simpiont/full-detailed-run.py](#) to validate our predicted performance using SMARTS.



03-SMARTS

Let's look at the `smarts_generator` with the statistical parameters.

- `n`: the number of the samples. It is a number of count.
- `k`: the systematic sampling interval. It is a number of count.
- `U`: the sampling unit size. It is a number of the instruction executed.
- `W`: the length of the detailed warm up period. It is a number of the instruction executed.

Each interval instruction length is $k*U$.

The warmup part starts at $(k-1)*U-W$

The detailed simulation part starts at $(k-1)*U$

This exit generator only works with SwitchableProcessor.

When it reaches to the start of the detailed warmup part, it resets the stats; then it switches the core type and schedule for the end of the warmup part and the end of the interval. When it reaches to the end of the detailed warmup part, it resets the stats. When it reaches to the end of the detailed simulation, it dumps the stats; then it switches the core type and schedule for the start of the next detailed warmup part.

03-SMARTS

In order to determine the k and w , we first need to determine what is our ideal n and ideal n . In the SMARTS paper, they define a large sample size as $n > 30$. For this exercise, let's set our ideal n as 50.

After setting up the ideal n , we can use it with the instructions executed in the workload to determine the ideal k .

Here is an example of how it can be done in python

```
program_length = 9115640
ideal_region_length = math.ceil(program_length/50)
ideal_U = 1000
ideal_k = math.ceil(ideal_region_length/ideal_U)
ideal_W = 2 * ideal_U
```

With this parameters, we will have about 50 samples. Each sample has a gap $(k-1)*U=182000$ instructions from each other. Each sample has 1000 instructions in detailed measurement and 2000 instructions of warm up period.



03-SMARTS

With these parameters, what the `smarts_generator` will

03-SMARTS

Now we understand what the `smarts_generator` will do, let's look at the system we are using on the [materials/02-Using-gem5/09-sampling/03-SMARTS/SMARTS.py](#).

It uses exactly the same system as the [materials/02-Using-gem5/09-sampling/01-simpoint/full-detailed-run.py](#), except that it uses a `SimpleSwitchableProcessor` to switch between ATOMIC CPU and O3 CPU for fast-forwarding and detailed simulation.

If the simulation is ran with

```
gem5 -re --outdir=full-detailed-run-m5out full-detailed-run.py
```

we can run the [materials/02-Using-gem5/09-sampling/03-SMARTS/predict_ipc.py](#) with to predict the overall IPC and calculate the relative error with the baseline IPC

```
python3 predict_ipc.py
```



03-SMARTS

What [materials/02-Using-gem5/09-sampling/03-SMARTS/predict_ipc.py](#) does is read the IPC from each sample and sum them up, then divided with the total number of samples to calculate the average IPC of the program.

$$\text{Average IPC} = \frac{\sum_{i=1}^n \text{IPC}_i}{n}$$

Here is what we expect to see

```
Number of samples: 50  
Predicted Overall IPC: 1.2563117400000001  
Actual Overall IPC: 1.247741  
Relative Error: 0.6869005667041583%
```

As the output suggests, the relative error between the predicted IPC and the actual baseline IPC is around 0.69%.

Summary

Now we experimented both targeted and statistical sampling in gem5, let's end it with the trade-offs.

Trade-off Method	Speedup	Resource Requirement	Prediction accuracy	Cost to change configuration	Ability to compare samples
Targeted Sampling	$\frac{S_{d/f}}{(S_{d/f} - 1) \times (n/N) + 2}$ $\frac{N}{n}$	Storage	It is impacted by the cluster signature and clustering algorithm	Almost no cost to change hardware configurations, but big cost to change software configurations	Allows directly comparison between samples
Statistical Sampling	$\frac{S_{d/f}}{(S_{d/f} - 1) \times (n/N) + 1}$	KVM & parallel resources	It is impacted by the input parameters	Equal cost to change both hardware and software configurations	Do not have comparable samples



Question in Mind

When we run the SimPoints, we need to run a separate simulation for each SimPoint, what if we can do it all in one simulation?